


Safety-relevant errors — F1 vs D



Day-out (leave-one-recording-day-out), 5 folds, 95 missions, seeds 42/123 (2) averaged per window. Shares from v29's pooled table (computed from the per-window predictions); raw counts are in the printed table. F1 and D are statistically indistinguishable (mission macro-F1 0.894 [0.832, 0.948] vs 0.893 [0.832, 0.948]). F1 is reported as the headline architecture because it has an interpretable, content-adaptive fusion mechanism tied to the project's complementarity evidence — not because it measurably outperforms D.