

Headline day-out results — F1 (headline) and D (comparison)

	accuracy	macro-F1 [95% CI]	balanced acc.	macro-AUROC
F1 — window	0.869	0.869 [0.821, 0.913]	0.869	0.973
F1 — mission	0.895	0.894 [0.832, 0.948]	0.895	0.986
D — window	0.872	0.871 [0.823, 0.917]	0.872	0.961
D — mission	0.895	0.893 [0.832, 0.948]	0.893	0.983

Day-out (leave-one-recording-day-out), 5 folds, 95 missions, seeds 42/123 (2) averaged per window. CI = mission-level cluster bootstrap (2000 resamples within class). F1 and D are statistically indistinguishable (mission macro-F1 0.894 [0.832, 0.948] vs 0.893 [0.832, 0.948]). F1 is reported as the headline architecture because it has an interpretable, content-adaptive fusion mechanism tied to the project's complementarity evidence — not because it measurably outperforms D.